Homework 5 Writeup

Instructions

- Provide an overview about how your project functions.
- Describe any interesting decisions you made to write your algorithm.
- Show and discuss the results of your algorithm.
- Feel free to include code snippets, images, and equations.
- List any extra credit implementation and result (optional).
- Use as many pages as you need, but err on the short side.
- **Please make this document anonymous.**

Project Overview

This project is about training a convolutional neural network to perform scene classification. We train a head on a pre-trained model, and create and train a model from scratch.

Implementation Detail

Result

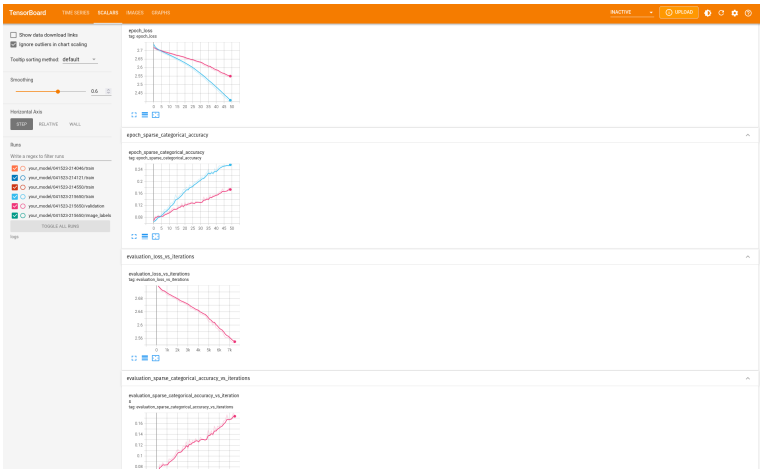

Figure 1: Task 1:Training the small model on CPU. Somehow this was very fast.

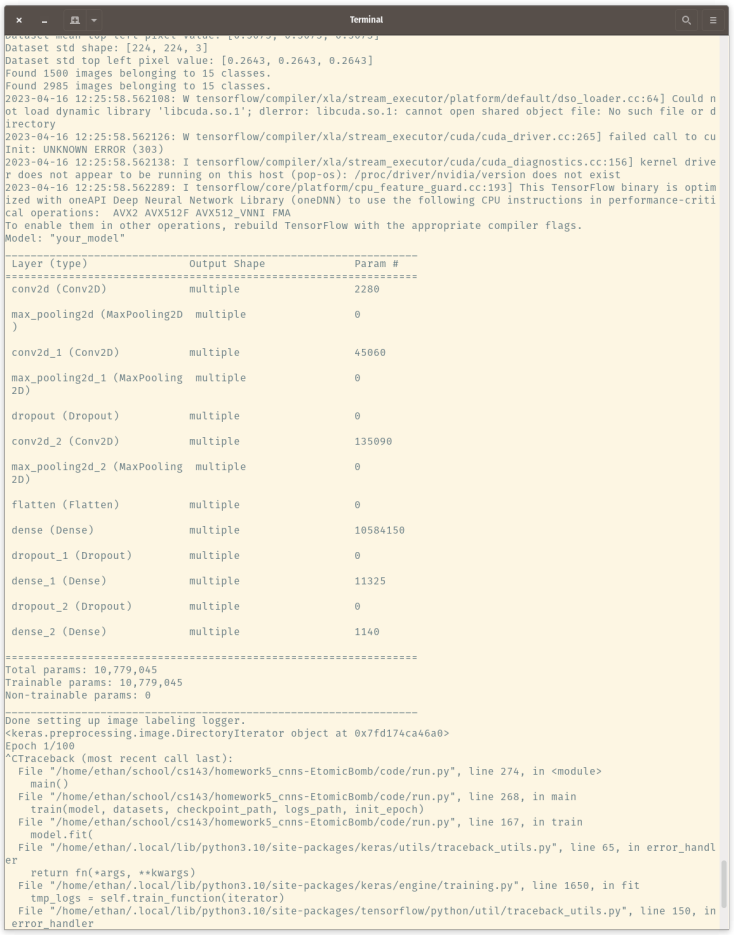

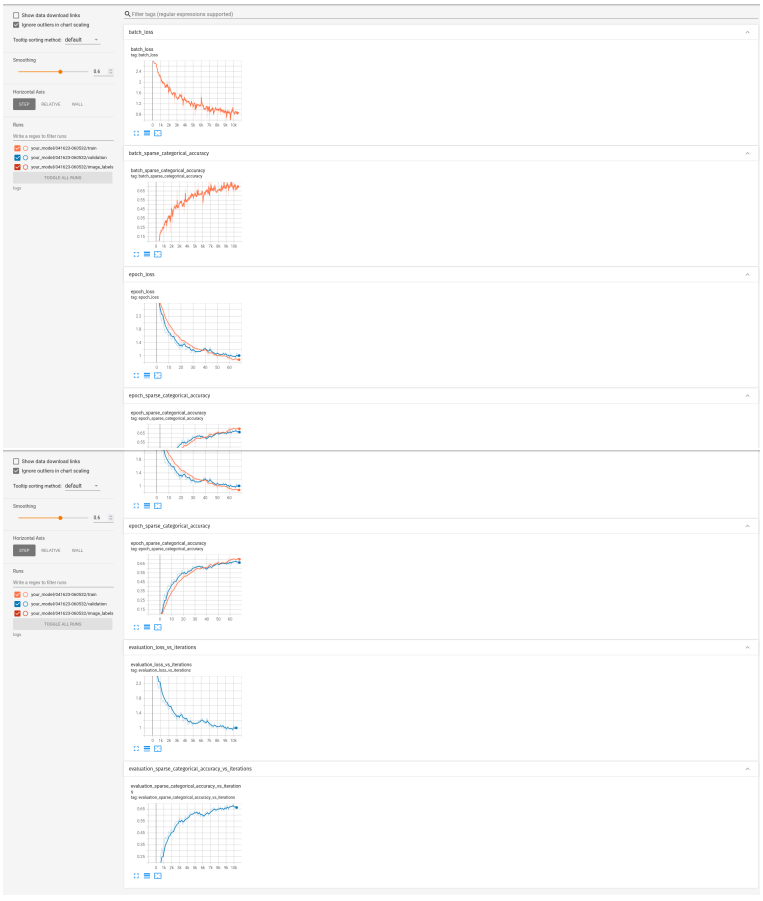

Figure 2: Progress training the complete model

With	Epochs	Sparse categorization accuracy
Base	3	0.2315
Standardization	3	0.1300
Data augmentation	3	0.1008
Regularization	56	0.6873

Table 1: Task 1 performance after adding each intermediate sophistication. I was able to train the network with full standardization, augmentation, and regularization on Colab. The intermediate networks were trained on my laptop (which only has a cpu – thus limited epochs). This is because I was not able to access Colab GPUs due to availability limits when I went back to do at the intermediate stages. I believe the decreasing accuracy as more sophistication was added is just because of random chance and high variation from few epochs. I noticed that before adding data augmentation (on my lost intermediate run, when I still could use Colab) the training accuracy was able to get to basically 100, while the test accuracy stayed low - maybe in the 20s or 30s. When I added augmentation, the accuracy in training and testing coincided.

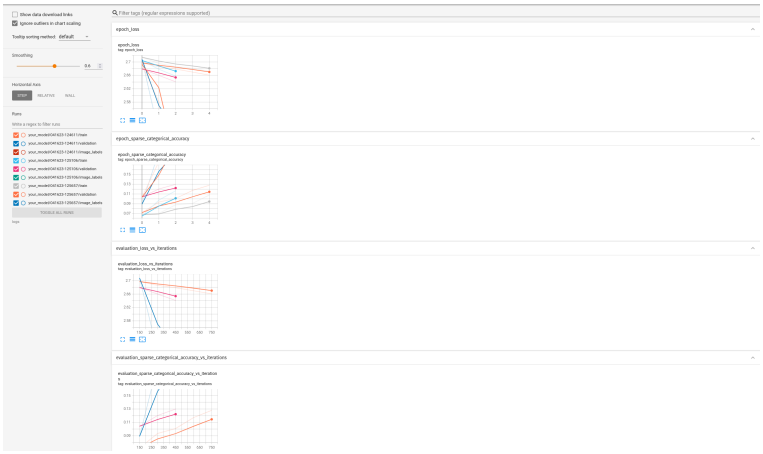

Figure 3: Task 1: Progress training the intermediate models.

Type	Sparse categorization accuracy
Total	0.9106

Table 2: Task 3 performance

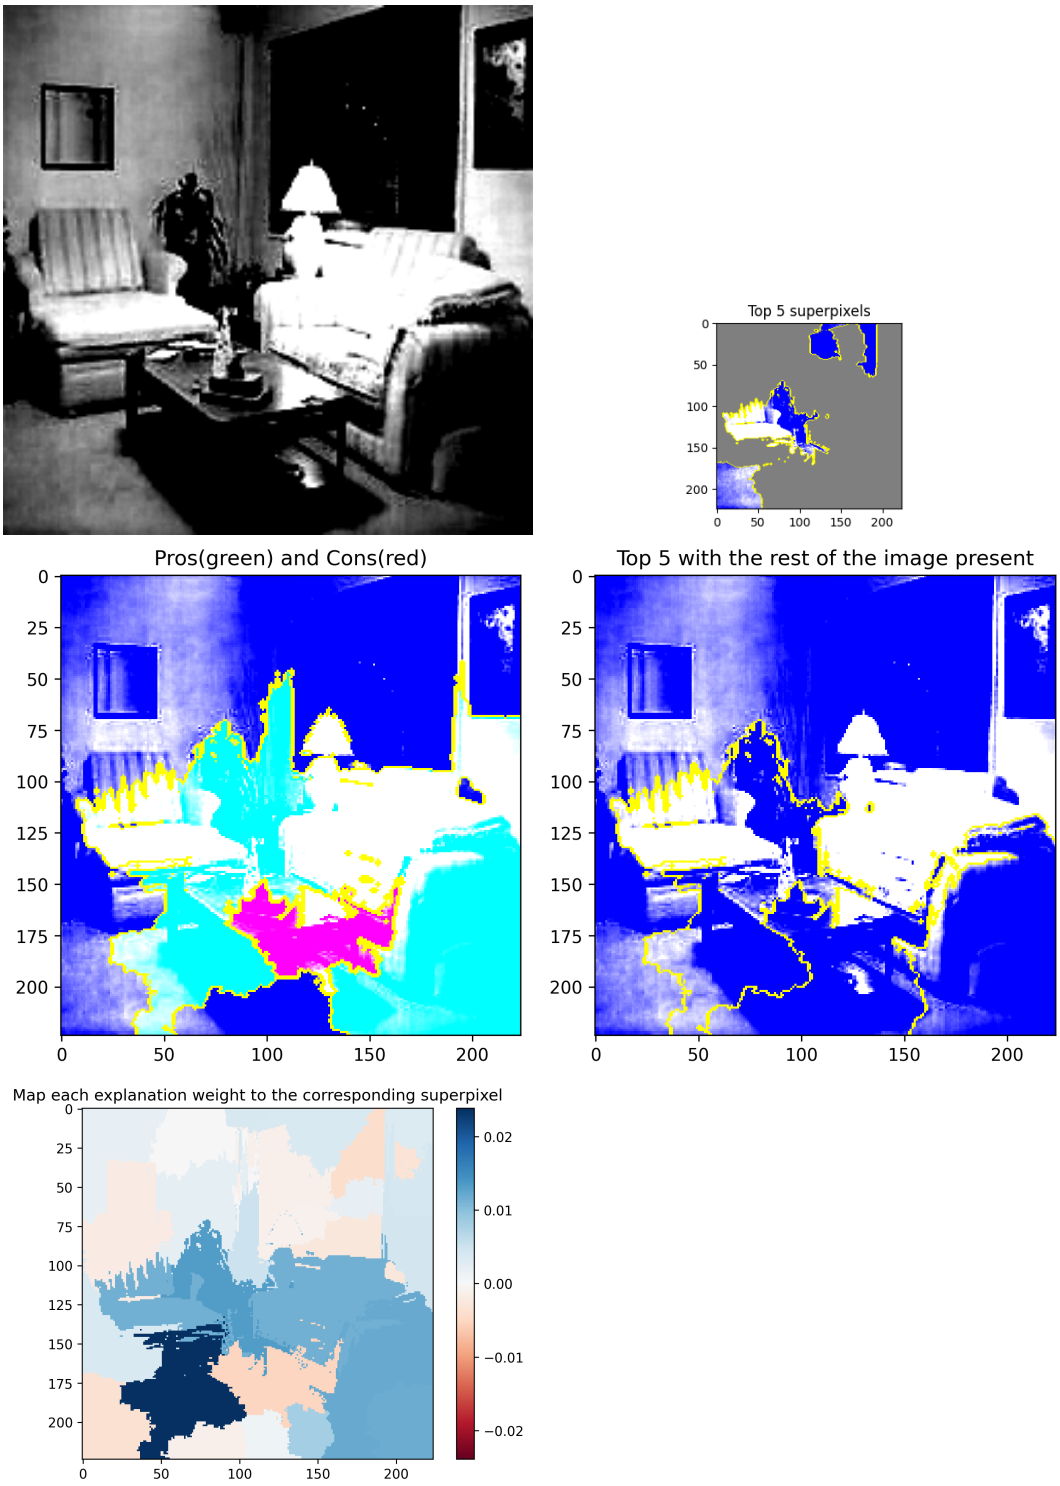

Figure 4: This is a bedroom, but was predicted as a living room. Curious. This model seemed to focus on the middle left of the image, especially the floor and part of the chair. It also seemed to attend to the window. These features – floors, chairs, and windows are all part of both bedrooms and living rooms. If it only attended to these features it's not suprising that it would get confused. I was not able to figure out whether it likes using the sky for its successful classifications.

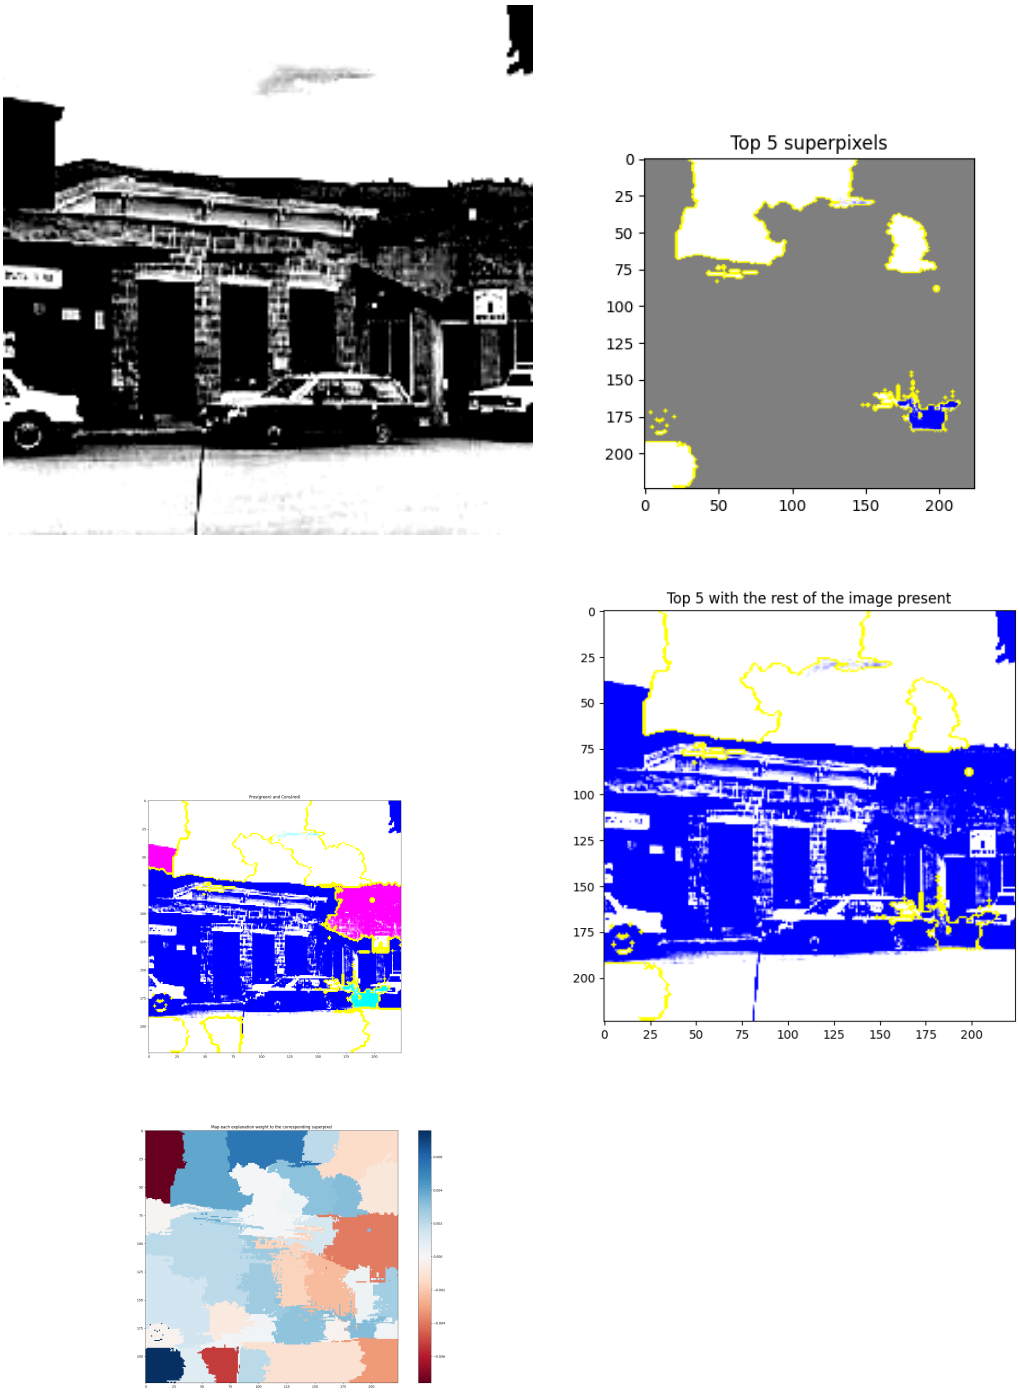

Figure 5: This is an inside city, but was predicted as a suburb. Curious. It seems like the model didn't really care about the architecture. Instead it focused on the sky and the front of a car, and part of the street. It would make sense that it would confuse an outdoor image of a suburb with a city if it was just looking at the sky.

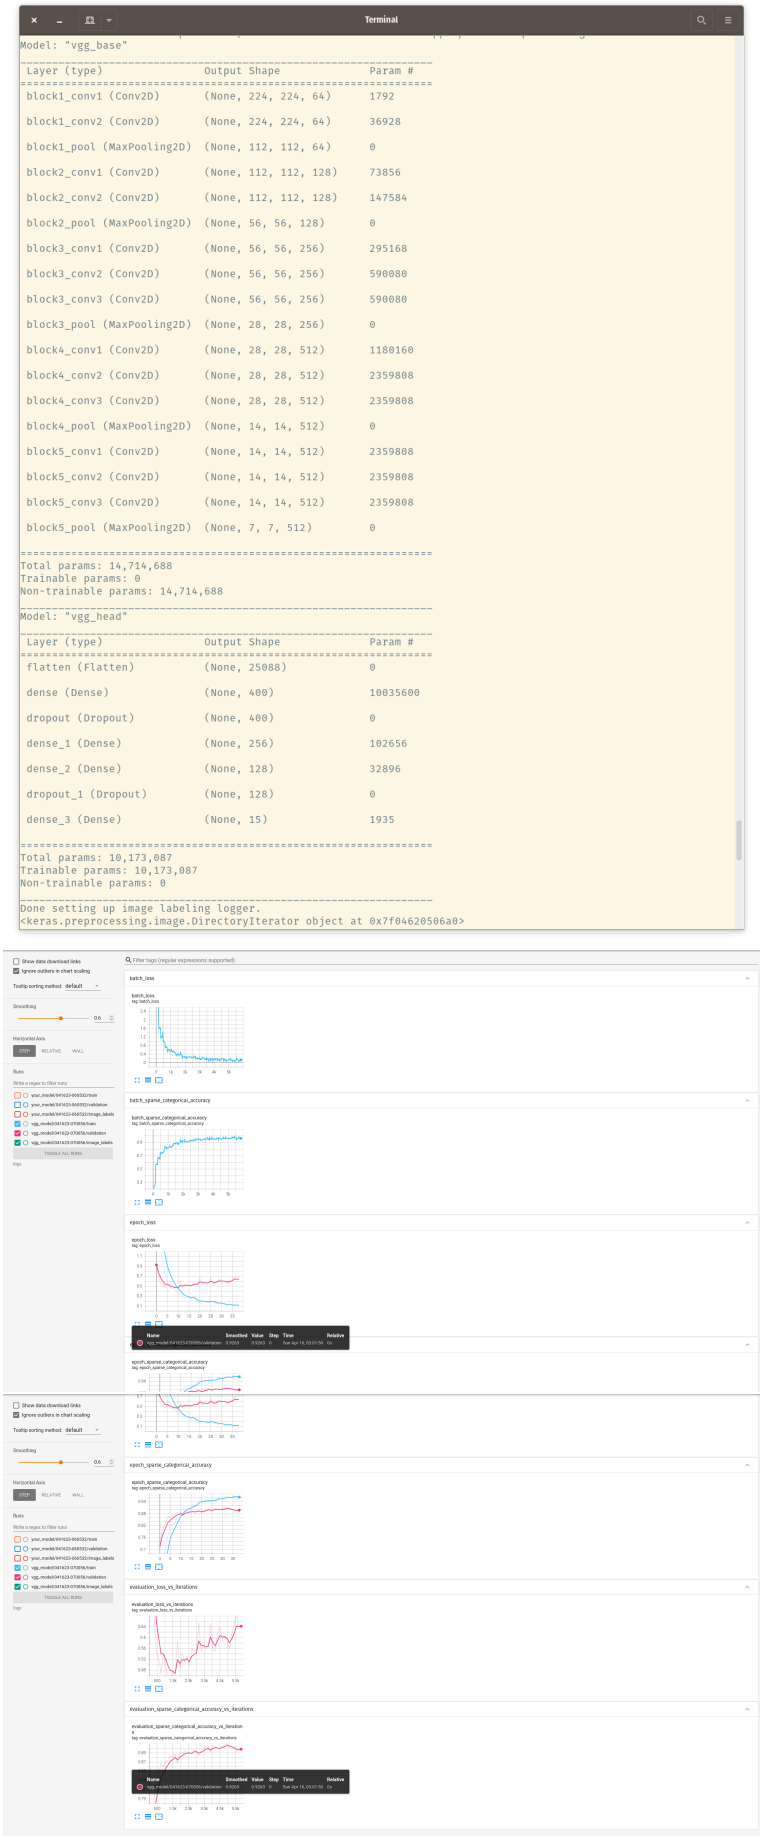

Figure 6: Task 3: Progress training head of the vgg model.